

# Problem Set 3: Window-Based Stereo Matching

## Description

In class and in Forsyth and Ponce, chapter 7 we discussed window-based approaches to estimating dense stereo correspondence. In this problem set you will implement such approaches and evaluate it on some standard stereo pairs.

## What to submit

Download and unzip: [ps3.zip](#)


ps3/

- input/ - input images, videos or other data supplied with the problem set
- output/ - directory containing output images and other files your code generates
- ps3.py - main code for completing each part, esp. function calls
- \*.py - Python modules, any utility code
- ps3\_report.pdf - a PDF file with all output images and text responses

Zip it as `ps3.zip`, and submit on T-Square.

## Questions

1. Use the pair [pair0-L.png](#) and [pair0-R.png](#) - a central square moved 2 pixels horizontally:



Implement the basic stereo algorithm of taking a window around every pixel in one image, and search for the best match along the same scan line in the other image. **You will do this both left to right and right to left.** Remember: Because of depth changes (discontinuities), some pixels visible in the left image are not in the right image and vice a versa. So you will match in both directions.

For this part, implement the simplest thing imaginable: Look for the smallest difference between the template window (source) and the proposed location window. Use the *sum of squared differences measure (SSD)*. We are going to take the definitions from: <https://software.intel.com/en-us/node/504333>

SSD is defined by:

$$S_{tx}(r, c) = \sum_{j=0}^{tplRows-1} \sum_{i=0}^{tplCols-1} \left[ t(j, i) - x\left(r+j-\frac{tplRows}{2}, c+i-\frac{tplCols}{2}\right) \right]^2$$

Basically, you just sum up the squares. A “good” match, then, is when this value is at a minimum. That is, you are looking for the same image patch in both images.

- a. Implement the SSD match algorithm as function **disparity\_ssd**(L, R) that returns a disparity image  $D(y, x)$  such that  $L(y, x) = R(y, x + D(y, x))$  when matching from left (L) to right (R).

The images supplied to this function (L and R) should be of the same size (height, width), single-channel (grayscale), and of double type with values in range  $[0.0, 1.0]$ . The disparity map returned should be of the same size as L and R, of double type, and each element should contain the disparity at that point, in pixels.

Apply it to the two test images, matching from left to right:

```
L = cv2.imread(... 'pair0-L.png' ...) * (1 / 255.0) # scale to [0, 1]
R = cv2.imread(... 'pair0-R.png' ...) * (1 / 255.0)
D_L = disparity_ssd(L, R)
```

Also match from right to left:

```
D_R = disparity_ssd(R, L)
```

They should indicate a central square moved 2 pixels to the left or right, e.g.,  $D_L$  should have value -2 in the approximate region of the central square, 0 elsewhere.

**Function:** `disparity_ssd`

**Output:** Save disparity images:

-  $D_L(y, x)$  [matching from left to right] as `ps3-1-a-1.png`

-  $D_R(y, x)$  [matching from right to left] as `ps3-1-a-2.png`

These disparity images may need to be scaled and shifted to display/write correctly.

2. Now we’re going to try this on a real image pair: `pair1-L.png` and `pair1-R.png`. Note that these are color images - so make sure you read in grayscale or convert.

- a. Again, apply your SSD match function, and create a disparity image  $D(y, x)$  such that  $L(y, x) = R(y, x + D(y, x))$  when matching from left to right. Also, match from right to left.

**Output:** Save disparity images, scaling/shifting as necessary:

-  $D_L(y, x)$  [matching from left to right] as `ps3-2-a-1.png`

-  $D_R(y, x)$  [matching from right to left] as `ps3-2-a-2.png`

- b. In the input directory are ground truth disparity images `pair1-D_L.png` and `pair1-D_R.png`. Compare your results.

**Output:** Text response - description of the differences between your results and ground truth.

3. SSD is not very robust to certain perturbations. We're going to try to see the effect of perturbations:
  - a. Using pair1, add some Gaussian noise, either to one image or both among pair1-L.png and pair1-R.png. Make the noise sigma big enough that you can tell some noise has been added. Run SSD match again.  
**Output:** Disparity images ( $D_L$  as ps3-3-a-1.png and  $D_R$  as ps3-3-a-2.png), text response - analysis of result compared to question 2.
  - b. Instead of the Gaussian noise, increase the contrast (multiplication) of one of the images by just 10%. Run SSD match again.  
**Output:** Disparity images ( $D_L$  as ps3-3-b-1.png and  $D_R$  as ps3-3-b-2.png), text response - analysis of result compared to question 2.
4. Now you're going to *use* (not implement yourself unless you want) an improved method, a similarity measure called **normalized correlation** – this is discussed in the book. The basic idea is that we think of two image patches as **vectors** and compute the angle between them – much like normalized dot products.

The explicit dot product of two image patches (treated as flat vectors) is:

$$R_{tx}(r, c) = \sum_{j=0}^{tplRows-1} \sum_{i=0}^{tplCols-1} t(j, i) \cdot x\left(r+j-\frac{tplRows}{2}, c+i-\frac{tplCols}{2}\right)$$

This result is then normalized:

$$\rho_{tx}(r, c) = \frac{R_{tx}(r, c)}{\sqrt{R_{xx}(r, c) R_{tt}\left(\frac{tplRows}{2}, \frac{tplCols}{2}\right)}}$$

- a. Using some form of normalized correlation, implement a window matching stereo algorithm. Again, write this as a function **disparity\_ncorr(L, R)** that returns a disparity image  $D(y, x)$  such that  $L(y, x) = R(y, x+D(y, x))$  when matching from left (L) to right (R). OpenCV has a variety of relevant functions and supported methods such as CV\_TM\_CCOEFF\_NORMED, which implement correlation similar to:

$$\gamma(u, v) = \frac{\sum_{x,y} [f(x, y) - \bar{f}_{u,v}] [t(x-u, y-v) - \bar{t}]}{\left\{ \sum_{x,y} [f(x, y) - \bar{f}_{u,v}]^2 \sum_{x,y} [t(x-u, y-v) - \bar{t}]^2 \right\}^{0.5}}$$

**You MAY use these built-in normalized correlation functions.**

Test it on the original images both left to right and right to left (pair1-L.png and pair1-R.png).

**Output:** Disparity images ( $D_L$  as ps3-4-a-1.png and  $D_R$  as ps3-4-a-2.png),

Text response - description of how it compares to the SSD version and to the ground truth.

- b. Now test it on both the noisy and contrast-boosted versions of pair1 from 3-a and 3-b.

**Output:** Disparity images (Gaussian noise:  $D_L$  as ps3-4-b-1.png and  $D_R$  as ps3-4-b-2.png; contrast-boosted:  $D_L$  as ps3-4-b-3.png and  $D_R$  as ps3-4-b-4.png),

Text response - analysis of results comparing original to noise and contrast-boosted images.

5. Finally, there is a second pair of images: pair2-L.png and pair2-R.png

- a. Try your algorithms on pair2. Play with the images – smooth, sharpen, etc. Keep comparing your results to the ground truth (pair2-D\_L.png and pair2-D\_R.png).

**Output:** Disparity images ( $D_L$  as ps3-5-a-1.png and  $D_R$  as ps3-5-a-2.png), Text response - analysis of what it takes to make stereo work using a window based approach.